

Binary Classification of Heart Disease of Patients using Deep Neural Network

Import Dataset

In [1]: `!pip install scikit-learn`

```

----- 2.1/9.2 MB 2.1 MB/s eta 0:00:04
----- 2.2/9.2 MB 2.1 MB/s eta 0:00:04
----- 2.3/9.2 MB 2.1 MB/s eta 0:00:04
----- 2.5/9.2 MB 2.1 MB/s eta 0:00:04
----- 2.6/9.2 MB 2.1 MB/s eta 0:00:04
----- 2.7/9.2 MB 2.1 MB/s eta 0:00:04
----- 2.8/9.2 MB 2.1 MB/s eta 0:00:04
----- 2.9/9.2 MB 2.2 MB/s eta 0:00:03
----- 3.0/9.2 MB 2.2 MB/s eta 0:00:03
----- 3.1/9.2 MB 2.2 MB/s eta 0:00:03
----- 3.2/9.2 MB 2.2 MB/s eta 0:00:03
----- 3.4/9.2 MB 2.2 MB/s eta 0:00:03
----- 3.4/9.2 MB 2.2 MB/s eta 0:00:03
----- 3.6/9.2 MB 2.2 MB/s eta 0:00:03
----- 3.7/9.2 MB 2.2 MB/s eta 0:00:03
----- 3.8/9.2 MB 2.2 MB/s eta 0:00:03
----- 3.9/9.2 MB 2.2 MB/s eta 0:00:03
----- 3.9/9.2 MB 2.2 MB/s eta 0:00:03
----- 4.0/9.2 MB 2.2 MB/s eta 0:00:03
----- 4.0/9.2 MB 2.2 MB/s eta 0:00:03

```

In [2]: `import pandas as pd`

In [3]: `df=pd.read_csv("heart_data.csv")`
`df.head()`

Out[3]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1

```
In [4]: df.shape
```

```
Out[4]: (303, 14)
```

```
In [5]: df.columns
```

```
Out[5]: Index(['age', 'sex', 'cp', 'trestbps', 'chol', 'fbs', 'restecg', 'thalach',  
              'exang', 'oldpeak', 'slope', 'ca', 'thal', 'target'],  
              dtype='object')
```

```
In [6]: df.size
```

```
Out[6]: 4242
```

Split the Dataset into Training and Testing

```
In [7]: x = df  
        y = df.pop('target')
```

```
In [8]: from sklearn.model_selection import train_test_split  
        x_train, x_test, y_train, y_test=train_test_split(x, y, test_size=0.2, random_state=42)
```

```
In [9]: x_train.shape
```

```
Out[9]: (242, 13)
```

```
In [10]: x_test.shape
```

```
Out[10]: (61, 13)
```

Create Neural Networks

```
In [11]: from tensorflow.keras.models import Sequential  
        from tensorflow.keras.layers import Dense
```

```
In [12]: model = Sequential()  
        model.add(Dense(8, input_dim=13, activation='relu'))  
        model.add(Dense(1, activation='sigmoid'))
```

Compile Model

```
In [13]: from tensorflow import keras
```

```
In [14]: optimizer = keras.optimizers.RMSprop(learning_rate=0.001)
```

```
In [15]: model.compile(loss='mse', optimizer=optimizer, metrics=['accuracy'])
model.fit(x_train, y_train, epochs=10, batch_size=30, verbose=1)
```

```
Epoch 1/10
9/9 [=====] - 3s 12ms/step - loss: 0.3713 - accuracy: 0.6074
Epoch 2/10
9/9 [=====] - 0s 1ms/step - loss: 0.3643 - accuracy: 0.6157
Epoch 3/10
9/9 [=====] - 0s 1ms/step - loss: 0.3609 - accuracy: 0.6240
Epoch 4/10
9/9 [=====] - 0s 1ms/step - loss: 0.3671 - accuracy: 0.6116
Epoch 5/10
9/9 [=====] - 0s 1ms/step - loss: 0.3692 - accuracy: 0.6033
Epoch 6/10
9/9 [=====] - 0s 1ms/step - loss: 0.3666 - accuracy: 0.6116
Epoch 7/10
9/9 [=====] - 0s 1ms/step - loss: 0.3647 - accuracy: 0.6033
Epoch 8/10
9/9 [=====] - 0s 1ms/step - loss: 0.3643 - accuracy: 0.6074
Epoch 9/10
9/9 [=====] - 0s 1ms/step - loss: 0.3603 - accuracy: 0.6116
Epoch 10/10
9/9 [=====] - 0s 1ms/step - loss: 0.3633 - accuracy: 0.6116
```

```
Out[15]: <keras.callbacks.History at 0x20059495480>
```

```
In [16]: model.evaluate(x_test,y_test)
```

```
2/2 [=====] - 0s 3ms/step - loss: 0.2826 - accuracy: 0.6885
```

```
Out[16]: [0.2826370596885681, 0.688524603843689]
```

Print the Model Summary

In [17]: `model.summary()`

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 8)	112
dense_1 (Dense)	(None, 1)	9
Total params: 121		
Trainable params: 121		
Non-trainable params: 0		

Train the Model

In [18]: `model.compile(loss='mse', optimizer=optimizer, metrics=['accuracy'])`
`model.fit(x_train, y_train, epochs=200, batch_size=10, verbose=1)`

```
Epoch 1/200
25/25 [=====] - 1s 1ms/step - loss: 0.3584 - accuracy: 0.6240
Epoch 2/200
25/25 [=====] - 0s 1ms/step - loss: 0.3736 - accuracy: 0.5950
Epoch 3/200
25/25 [=====] - 0s 1ms/step - loss: 0.3738 - accuracy: 0.6033
Epoch 4/200
25/25 [=====] - 0s 1ms/step - loss: 0.3717 - accuracy: 0.6074
Epoch 5/200
25/25 [=====] - 0s 1ms/step - loss: 0.3724 - accuracy: 0.5950
Epoch 6/200
25/25 [=====] - 0s 1ms/step - loss: 0.3601 - accuracy: 0.6116
Epoch 7/200
25/25 [=====] - 0s 1ms/step - loss: 0.3601 - accuracy: 0.6116
```

In [19]: `model.evaluate(x_test, y_test)`

```
2/2 [=====] - 0s 3ms/step - loss: 0.1279 - accuracy: 0.8525
```

Out[19]: [0.12794315814971924, 0.8524590134620667]

Save the Model and Split the data for Validation

```
In [20]: history = model.fit(x_train, y_train, validation_split=0.2, epochs=100, batch_size=32)
y: 0.8497 - val_loss: 0.1310 - val_accuracy: 0.8163
Epoch 4/100
20/20 [=====] - 0s 3ms/step - loss: 0.1194 - accuracy: 0.8446 - val_loss: 0.1393 - val_accuracy: 0.8163
Epoch 5/100
20/20 [=====] - 0s 3ms/step - loss: 0.1279 - accuracy: 0.8290 - val_loss: 0.1183 - val_accuracy: 0.8367
Epoch 6/100
20/20 [=====] - 0s 3ms/step - loss: 0.1235 - accuracy: 0.8342 - val_loss: 0.1066 - val_accuracy: 0.8571
Epoch 7/100
20/20 [=====] - 0s 3ms/step - loss: 0.1285 - accuracy: 0.8394 - val_loss: 0.1028 - val_accuracy: 0.8367
Epoch 8/100
20/20 [=====] - 0s 5ms/step - loss: 0.1155 - accuracy: 0.8394 - val_loss: 0.1272 - val_accuracy: 0.8163
Epoch 9/100
20/20 [=====] - 0s 4ms/step - loss: 0.1275 - accuracy: 0.8031 - val_loss: 0.1211 - val_accuracy: 0.8163
Epoch 10/100
-----
```

```
In [21]: model.evaluate(x_test,y_test)

2/2 [=====] - 0s 2ms/step - loss: 0.1432 - accuracy: 0.8361
```

```
Out[21]: [0.143158420920372, 0.8360655903816223]
```

Print the model accuracy


```
In [22]: history.history.keys()
```


```
Out[22]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

```
In [23]: import matplotlib.pyplot as plt
```

Matplotlib is building the font cache; this may take a moment.

```
In [24]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['Train', 'Validation'])
plt.show()
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['Train', 'Validation'])
plt.show()
```





Do further experiments

```
In [25]: model1 = Sequential()

model1.add(Dense(16, input_dim=13, activation='relu'))
model1.add(Dense(8, activation='relu'))
model1.add(Dense(1, activation='sigmoid'))
```

```
In [28]: model1.compile(loss='mse', optimizer=optimizer, metrics=['accuracy'])
model1.fit(x_train, y_train, epochs=10, batch_size=30, verbose=1)
```

```
Epoch 1/10
9/9 [=====] - 1s 2ms/step - loss: 0.4587 - accuracy: 0.5413
Epoch 2/10
9/9 [=====] - 0s 1ms/step - loss: 0.4587 - accuracy: 0.5413
Epoch 3/10
9/9 [=====] - 0s 1ms/step - loss: 0.4587 - accuracy: 0.5413
Epoch 4/10
9/9 [=====] - 0s 2ms/step - loss: 0.4587 - accuracy: 0.5413
Epoch 5/10
9/9 [=====] - 0s 1ms/step - loss: 0.4587 - accuracy: 0.5413
Epoch 6/10
9/9 [=====] - 0s 1ms/step - loss: 0.4587 - accuracy: 0.5413
Epoch 7/10
9/9 [=====] - 0s 2ms/step - loss: 0.4587 - accuracy: 0.5413
Epoch 8/10
9/9 [=====] - 0s 1ms/step - loss: 0.4587 - accuracy: 0.5413
Epoch 9/10
9/9 [=====] - 0s 2ms/step - loss: 0.4587 - accuracy: 0.5413
Epoch 10/10
9/9 [=====] - 0s 1ms/step - loss: 0.4587 - accuracy: 0.5413
```

```
Out[28]: <keras.callbacks.History at 0x2005e712b90>
```

```
In [30]: model1.evaluate(x_test, y_test)
```

```
2/2 [=====] - 0s 3ms/step - loss: 0.4426 - accuracy: 0.5574
```

```
Out[30]: [0.44262295961380005, 0.5573770403862]
```


In [32]: `history1 = model.fit(x_train, y_train, validation_split=0.2, epochs=100, batch_si`

```
Epoch 1/100
10/10 [=====] - 0s 10ms/step - loss: 0.1230 - accuracy: 0.8342 - val_loss: 0.1074 - val_accuracy: 0.8776
Epoch 2/100
10/10 [=====] - 0s 5ms/step - loss: 0.1151 - accuracy: 0.8446 - val_loss: 0.1190 - val_accuracy: 0.8571
Epoch 3/100
10/10 [=====] - 0s 5ms/step - loss: 0.1243 - accuracy: 0.8290 - val_loss: 0.1185 - val_accuracy: 0.8163
Epoch 4/100
10/10 [=====] - 0s 4ms/step - loss: 0.1152 - accuracy: 0.8394 - val_loss: 0.1163 - val_accuracy: 0.8571
Epoch 5/100
10/10 [=====] - 0s 5ms/step - loss: 0.1244 - accuracy: 0.8187 - val_loss: 0.1069 - val_accuracy: 0.8776
Epoch 6/100
10/10 [=====] - 0s 4ms/step - loss: 0.1173 - accuracy: 0.8446 - val_loss: 0.1076 - val_accuracy: 0.8776
Epoch 7/100
10/10 [=====] - 0s 4ms/step - loss: 0.1244 - accuracy: 0.8187 - val_loss: 0.1069 - val_accuracy: 0.8776
```

In [33]: `model1.summary()`

Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 16)	224
dense_3 (Dense)	(None, 8)	136
dense_4 (Dense)	(None, 1)	9
Total params: 369		
Trainable params: 369		
Non-trainable params: 0		

In [34]: `ls = history1.history`

```
In [35]: new = pd.DataFrame.from_dict(ls)
new
```

Out[35]:

	loss	accuracy	val_loss	val_accuracy
0	0.122989	0.834197	0.107391	0.877551
1	0.115098	0.844560	0.118961	0.857143
2	0.124288	0.829016	0.118453	0.816327
3	0.115204	0.839378	0.116296	0.857143
4	0.124392	0.818653	0.106862	0.877551
...
95	0.117606	0.829016	0.106387	0.857143
96	0.119126	0.865285	0.124048	0.816327
97	0.122252	0.813471	0.117347	0.857143
98	0.113644	0.829016	0.107408	0.836735
99	0.115921	0.839378	0.106518	0.877551

100 rows × 4 columns

```
In [36]: model2 = Sequential()
model2.add(Dense(32, input_dim=13, activation='relu'))
model2.add(Dense(16, activation='relu'))
model2.add(Dense(8, activation='relu'))
model2.add(Dense(1, activation='sigmoid'))
```

```
In [38]: model2.compile(loss='mse', optimizer=optimizer, metrics=['accuracy'])  
model2.fit(x_train, y_train, epochs=10, batch_size=30, verbose=1)
```

```
Epoch 1/10  
9/9 [=====] - 1s 2ms/step - loss: 0.4433 - accuracy:  
0.4669  
Epoch 2/10  
9/9 [=====] - 0s 2ms/step - loss: 0.3428 - accuracy:  
0.5372  
Epoch 3/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2955 - accuracy:  
0.5455  
Epoch 4/10  
9/9 [=====] - 0s 1ms/step - loss: 0.3113 - accuracy:  
0.5661  
Epoch 5/10  
9/9 [=====] - 0s 1ms/step - loss: 0.2619 - accuracy:  
0.5992  
Epoch 6/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2694 - accuracy:  
0.6116  
Epoch 7/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2646 - accuracy:  
0.5785  
Epoch 8/10  
9/9 [=====] - 0s 1ms/step - loss: 0.2771 - accuracy:  
0.6157  
Epoch 9/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2367 - accuracy:  
0.6157  
Epoch 10/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2324 - accuracy:  
0.6405
```

```
Out[38]: <keras.callbacks.History at 0x2005fd6df60>
```

```
In [40]: model2.evaluate(x_test, y_test)
```

```
2/2 [=====] - 0s 4ms/step - loss: 0.2237 - accuracy:  
0.6721
```

```
Out[40]: [0.22372157871723175, 0.6721311211585999]
```

```
In [41]: model2.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 32)	448
dense_6 (Dense)	(None, 16)	528
dense_7 (Dense)	(None, 8)	136
dense_8 (Dense)	(None, 1)	9

=====
Total params: 1,121

Trainable params: 1,121

Non-trainable params: 0
=====

```
In [42]: model3 = Sequential()  
model3.add(Dense(64, input_dim=13, activation='relu'))  
model3.add(Dense(32, activation='relu'))  
model3.add(Dense(16, activation='relu'))  
model3.add(Dense(8, activation='relu'))  
model3.add(Dense(1, activation='sigmoid'))
```

```
In [44]: model3.compile(loss='mse', optimizer=optimizer, metrics=['accuracy'])  
model3.fit(x_train, y_train, epochs=10, batch_size=30, verbose=1)
```

```
Epoch 1/10  
9/9 [=====] - 1s 2ms/step - loss: 0.4796 - accuracy:  
0.4876  
Epoch 2/10  
9/9 [=====] - 0s 2ms/step - loss: 0.3466 - accuracy:  
0.5661  
Epoch 3/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2638 - accuracy:  
0.5702  
Epoch 4/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2510 - accuracy:  
0.5744  
Epoch 5/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2530 - accuracy:  
0.5868  
Epoch 6/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2355 - accuracy:  
0.6240  
Epoch 7/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2056 - accuracy:  
0.7025  
Epoch 8/10  
9/9 [=====] - 0s 3ms/step - loss: 0.2299 - accuracy:  
0.6488  
Epoch 9/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2102 - accuracy:  
0.6736  
Epoch 10/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2239 - accuracy:  
0.6529
```

```
Out[44]: <keras.callbacks.History at 0x2005e302320>
```

```
In [46]: model3.evaluate(x_test, y_test)
```

```
2/2 [=====] - 0s 3ms/step - loss: 0.2123 - accuracy:  
0.6885
```

```
Out[46]: [0.2122710943222046, 0.688524603843689]
```

In [47]: `model3.summary()`

Model: "sequential_3"

Layer (type)	Output Shape	Param #
=====		
dense_9 (Dense)	(None, 64)	896
dense_10 (Dense)	(None, 32)	2080
dense_11 (Dense)	(None, 16)	528
dense_12 (Dense)	(None, 8)	136
dense_13 (Dense)	(None, 1)	9
=====		
Total params: 3,649		
Trainable params: 3,649		
Non-trainable params: 0		

In []: